

HW2

Que1.1) C code attached as hw2_a1.c

```
#include "mpi.h"
#include <stdlib.h>
#include <string.h>
#include <time.h>
#include <math.h>

#define pi 3.14159265358979323846 /* pi */

/*MPI_Type Enum*/
enum {
    MASTER = 0,
    FROM_MASTER = 1,
    FROM_WORKER = 2,
};

double estimate_g(double lower_bound, double upper_bound, long long int N)
{
    int          num_tasks;      /* number of tasks */
    int          task_id;        /* number of processes */
    int          num_workers;     /* number of worker tasks */
    int          source;          /* rank of sender */
    int          dest;            /* rank of receiver */
    int          mtype;           /* message type */
    MPI_Status    status;         /* return status for receive */
    int          rows;
    double        x, y;           /* First boundary condition, Second boundary
condition */
    double        sum1 = 0, sum = 0;
    int          i;
    clock_t begin, end;
    double time_spent;

    /* Find out process rank */
    MPI_Comm_rank(MPI_COMM_WORLD, &task_id);

    /* Find out number of processes */
    MPI_Comm_size(MPI_COMM_WORLD, &num_tasks);

    num_workers = num_tasks - 1;
    int average_row = N / num_workers;
    int left_overs = N % num_workers;
```

```

        if (task_id == MASTER) {
            //begin = clock();
            //Send matrix data to worker tasks
            mtype = FROM_MASTER;
            for (dest = 1; dest < num_tasks; dest++) {
                rows = (dest <= left_overs) ? average_row + 1 : average_row;
                printf("Sending %d rows to task %d\n", rows, dest);
                MPI_Send(&lower_bound, 1, MPI_DOUBLE, dest, mtype,
MPI_COMM_WORLD);
                MPI_Send(&upper_bound, 1, MPI_DOUBLE, dest, mtype,
MPI_COMM_WORLD);
                MPI_Send(&N, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
                MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
            }
            printf("Sent all the data!\n");
        }
        else {
            //Worker task
            mtype = FROM_MASTER;

            MPI_Recv(&lower_bound, 1, MPI_DOUBLE, MASTER, mtype, MPI_COMM_WORLD,
&status);
            MPI_Recv(&upper_bound, 1, MPI_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);
            MPI_Recv(&N, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
            MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
            double result, final_val = 0, sum = 0;
            int j;
            srand(time(NULL)*task_id);
            for (j = 0; j < rows; j++)
            {
                result = exp(-1 * pow(2 * ((rand() / (RAND_MAX / (upper_bound -
lower_bound)))) + lower_bound), 2));
                final_val += result;
            }
            sum1 = (8 * sqrt(2 * pi))*((upper_bound - lower_bound) / N)*final_val;
        }
        return sum1;
    }

void collect_results(double *result)
{
    int          num_tasks;          /* number of tasks */
    int          task_id;             /* number of processes */
    int          source;               /* rank of sender */
    int          dest;                 /* rank of receiver */

```

```

int          mtype;          /* message type */
MPI_Status   status;         /* return status for receive */
double       sum1 = 0, sum = 0;
int          i;

/* Find out number of processes */
MPI_Comm_size(MPI_COMM_WORLD, &num_tasks);

/* Find out process rank */
MPI_Comm_rank(MPI_COMM_WORLD, &task_id);

sum1 = *result;

//Wait for results from workers
if (task_id == MASTER)
{
    mtype = FROM_WORKER;
    for (i = 1; i < num_tasks; i++) {
        source = i;
        printf("Receiving data from task %d\n", source);
        MPI_Recv(&sum1, 1, MPI_DOUBLE, source, mtype,
MPI_COMM_WORLD, &status);
        sum += sum1;
    }
    //Print the results
    printf("\n\nValue of integration is:%lf\n", sum);
}
else
{
    //Send the results back to the Master process
    mtype = FROM_WORKER;
    MPI_Send(&sum1, 1, MPI_DOUBLE, MASTER, mtype, MPI_COMM_WORLD);
}
}

int main(int argc, char *argv[]) {

    int          num_tasks;      /* number of tasks */
    int          task_id;        /* number of processes */
    clock_t begin, end;
    double time_spent;

    //MPI
    /* Start up MPI */
    MPI_Init(&argc, &argv);

```

```

/* Find out number of processes */
MPI_Comm_size(MPI_COMM_WORLD, &num_tasks);

/* Find out process rank */
MPI_Comm_rank(MPI_COMM_WORLD, &task_id);

float lower_bound = atof(argv[1]);
float upper_bound = atof(argv[2]);
long long int N = atof(argv[3]);

begin = clock();

double result = estimate_g(lower_bound, upper_bound, N);
collect_results(&result);

MPI_Finalize();
if (task_id == MASTER) {

    end = clock();
    time_spent = (double)(end - begin) / CLOCKS_PER_SEC;
    printf("%d number of processors and time spent %f sec\n", num_tasks,
time_spent);
}
return 0;
}

```

Que1.2) C code attached as hw2_a2.c

```

#include "mpi.h"
#include <stdlib.h>
#include <string.h>
#include <time.h>
#include <math.h>

```

```

#define pi 3.14159265358979323846 /* pi */

```

```

/*MPI_Type Enum*/

```

```

enum {
    MASTER = 0,
    FROM_MASTER = 1,
    FROM_WORKER = 2,
};

```

```

double estimate_g(double lower_bound, double upper_bound, long long int N)

```

```

{
    int          num_tasks;          /* number of tasks */
    int          task_id;            /* number of processes */
    int          num_workers;        /* number of worker tasks */

```

```

int          source;          /* rank of sender */
int          dest;            /* rank of receiver */
int          mtype;           /* message type */
MPI_Status   status;          /* return status for receive */
int          rows;
double       x, y;            /* First boundary condition, Second boundary
condition */
double       sum1 = 0, sum = 0;
int          i;
clock_t      begin, end;
double       time_spent;

/* Find out process rank */
MPI_Comm_rank(MPI_COMM_WORLD, &task_id);

/* Find out number of processes */
MPI_Comm_size(MPI_COMM_WORLD, &num_tasks);

num_workers = num_tasks - 1;
int average_row = N / num_workers;
int left_overs = N % num_workers;

if (task_id == MASTER) {
    //begin = clock();
    //Send matrix data to worker tasks
    mtype = FROM_MASTER;
    for (dest = 1; dest < num_tasks; dest++) {
        rows = (dest <= left_overs) ? average_row + 1 : average_row;
        printf("Sending %d rows to task %d\n", rows, dest);
        MPI_Send(&lower_bound, 1, MPI_DOUBLE, dest, mtype,
MPI_COMM_WORLD);
        MPI_Send(&upper_bound, 1, MPI_DOUBLE, dest, mtype,
MPI_COMM_WORLD);
        MPI_Send(&N, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
        MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    }
    printf("Sent all the data!\n");
}
else {
    //Worker task
    mtype = FROM_MASTER;

    MPI_Recv(&lower_bound, 1, MPI_DOUBLE, MASTER, mtype, MPI_COMM_WORLD,
&status);

```

```

        MPI_Recv(&upper_bound, 1, MPI_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);
        MPI_Recv(&N, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
        MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
        double result, final_val = 0, sum = 0;
        int j;
        srand(time(NULL)*task_id);
        for (j = 0; j < rows; j++)
        {
            result = exp(-1 * pow(2 * ((rand() / (RAND_MAX / (upper_bound -
lower_bound)))) + lower_bound), 2));
            final_val += result;
        }
        sum1 = (8 * sqrt(2 * pi))*((upper_bound - lower_bound) / N)*final_val;
    }
    return sum1;
}

void collect_results(double *result)
{
    int            num_tasks;           /* number of tasks */
    int            task_id;             /* number of processes */

    double         sum1 = 0, sum = 0;
    int            i;

    /* Find out number of processes */
    MPI_Comm_size(MPI_COMM_WORLD, &num_tasks);

    /* Find out process rank */
    MPI_Comm_rank(MPI_COMM_WORLD, &task_id);

    sum1 = *result;

    //Wait for results from workers
    MPI_Reduce(&sum1, &sum, num_tasks, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
    if (task_id == MASTER)
    {
        printf("\n\n HERE Value of integration is:%f\n", sum);
    }
}

int main(int argc, char *argv[]) {
    int            num_tasks;           /* number of tasks */
    int            task_id;             /* number of processes */
    clock_t begin, end;
    double time_spent;

```

```

//MPI
/* Start up MPI */
MPI_Init(&argc, &argv);

/* Find out number of processes */
MPI_Comm_size(MPI_COMM_WORLD, &num_tasks);

/* Find out process rank */
MPI_Comm_rank(MPI_COMM_WORLD, &task_id);

float lower_bound = atof(argv[1]);
float upper_bound = atof(argv[2]);
long long int N = atof(argv[3]);

begin = clock();

double result = estimate_g(lower_bound, upper_bound, N);
collect_results(&result);

MPI_Finalize();

if (task_id == MASTER) {

    end = clock();
    time_spent = (double)(end - begin) / CLOCKS_PER_SEC;
    printf("time spent %f sec\n", time_spent);
}
return 0;
}
Que1.3)

```

Ans:

Code 1.1) Tests, Results and Timing information:

Lower bound is 0 and upper bound is 0.06

Case 1: For processor cores 32 and samples 10000000
 Value of integration is: 1.197431
 32 number of processors and time spent 0.040000 sec.

Case 2: For processor cores 8 and samples 1000
 Value of integration is: 1.197294
 8 number of processors and time spent 0.000000 sec

Case 3: For processor cores 4 and samples 10000000

Value of integration is: 1.197434
4 number of processors and time spent 0.250000 sec

Case 4: For processor cores 4 and samples 100
Value of integration is: 1.197434
4 number of processors and time spent 0.000000 sec

Code 1.2) Tests, Results and Timing information:

Lower bound is 0 and upper bound is 0.06

Case 1: For processor cores 32 and samples 10000000
Value of integration is: 1.197432
32 number of processors and time spent 0.020000 sec.

Case 2: For processor cores 8 and samples 1000
HERE Value of integration is: 1.197336
8 number of processors and time spent 0.000000 sec

Case 3: For processor cores 4 and samples 10000000
Value of integration is: 1.197431
4 number of processors and time spent 0.200000 sec

Case 4: For processor cores 4 and samples 100
Value of integration is: 1.197394
4 number of processors and time spent 0.000000 sec

As calculation for million sample on 32 processor cores took far less time than same number of calculations on 4 processor cores as evident with case 1 and 3 in both codes results.

The Code was embarrassingly parallel as all parts of the code could be parallelized without much effort as there was no interdependency between any tasks or results.

Que2.1) C code attached as hw2_a2.c

/******Source Code *****/

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>
// CPU Calculation to check the output
#define MAX_RANGE 10001
#define MASTER 0
```



```
void out_results(int *tst, int Nw, int Nh)
```

```
{
    int i, j;
    for (i = 0; i < Nh; i++)
    {
        for (j = 0; j < Nw; j++)
            printf("%d ", tst[i*Nw + j]);
        printf("\n");
    }
}
```

```
void initialize_data(int **A, int **Ap, int N)
```

```
{
    int i, j, p, rank;
    srand(time(NULL));
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    /*Master node will have full memory allocation*/
    //if (rank == 0)
    {
        *A = (int *)malloc(N * N * sizeof(int));
        *Ap = (int *)calloc(N * N, sizeof(int));
    }

    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            (*A)[i + N*j] = rand()% MAX_RANGE;
/*
    if (rank == 0)
    {
        printf("input matrix \n");
        out_results(*A, N, N);
    }
*/
}
```

```
void scatter_data(int *A, int N)
```

```
{
    int i, p, rank, compute_size;
    int *offsets, *chunk_sizes;
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    offsets = (int *)malloc(p*sizeof(int)); /*offsets from which data is sent of input A from
master*/
    chunk_sizes = (int *)malloc(p*sizeof(int)); /*sizes of the data chunks sent of input A from
master*/

    compute_size = ((N - 2) / p);
    for (i = 0; i < p; ++i) {
```

```

        /*data chunks that would be computed at each node*/
        chunk_sizes[i] = (compute_size + 2)*N + (i==0)*((N-2)%p)*N;
        offsets[i] = i*compute_size*N+(i!=0)*((N-2)%p)*N;
    }
    if (((N - 2) / p) > 0)
        /*Scatter data from the master to the worker*/
        MPI_Scatterv(A, chunk_sizes, offsets, MPI_INT, A, (compute_size+2)*N, MPI_INT,
                    MASTER, MPI_COMM_WORLD);

    free(chunk_sizes);
    free(offsets);
}

void mask_operation(int *A, int N, int * Ap)
{
    int i, j, p, rank, compute_size;
    int *res ;
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    /*Values converted to float before computation to avoid accumulation from overflowing as
    MAX_RANGE 10000 */
    /*Offset added to avoid first row overwrite for master*/
    if (rank == 0)
    {
        compute_size = ((N - 2) / p)+((N-2)%p);/*data chunks + leftovers that would be
        computed at each node*/
        res = Ap +N;
    }
    else
    {
        compute_size = ((N - 2) / p);/*data chunks that would be computed at each node*/
        res = Ap;
    }
    for (i = 1; i < compute_size+1; i++)
    {
        for (j = 1; j < N - 1; j++)
        {
            res[(i-1)*N + j] =(int)( ((float) A[(i - 1) * N + j - 1] + (float)A[(i - 1)*N + j] +
            (float)A[(i - 1)*N + j + 1] + (float)A[i *N + j - 1] +
            2 * ((float)A[i*N + j]) + (float)A[i*N + j + 1] + (float)A[(i + 1)*N
            + j - 1] + (float)A[(i + 1)*N + j] + (float)A[(i + 1)*N + j + 1]) / 10);
        }
    }
}

void gather_results(int *Ap, int N) {
    int i, p, rank, compute_size;
    int *offsets, *chunk_sizes;

```

```

int *res;
MPI_Comm_size(MPI_COMM_WORLD, &p);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
offsets = (int *)malloc(p*sizeof(int));/*offsets from which data is sent of input A from
master*/
chunk_sizes = (int *)malloc(p*sizeof(int));
compute_size = ((N - 2) / p);
for (i = 0; i < p; ++i) {
    /*data chunks that would be computed at each node*/
    chunk_sizes[i] = (compute_size + 2)*N + (i==0)*((N-2)%p)*N;
    offsets[i] = i*compute_size*N+(i!=0)*((N-2)%p)*N+N;
}
/*Offset added to avoid first row overwrite for master*/
if (rank == 0)
    res = Ap + N;
else
    res = Ap;
// Gather the generated outputs
if (((N - 2) / p) > 0)
    MPI_Gatherv(res, compute_size*N, MPI_INT, Ap, chunk_sizes, offsets, MPI_INT,
        MASTER, MPI_COMM_WORLD);
/*
if (rank == 0)
{
    printf("Output Matrix\n");
    out_results(Ap, N,N);
}*/
free(chunk_sizes);
free(offsets);
}
int main(int argc, char** argv) {

    int * A, *Ap,*tst;
    int p, rank;
    clock_t begin,end;
    double time_spent;
    MPI_Init(&argc, &argv);
    int N = atoi(argv[1]);
    begin = clock();
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    initialize_data(&A, &Ap,N);
    scatter_data(A, N);
    mask_operation(A, N, Ap);
    gather_results(Ap, N);
    if (rank == MASTER) {

        end = clock();
        time_spent = (double)(end - begin) / CLOCKS_PER_SEC;
    }
}

```

```

        printf("time spent %f sec\n", time_spent);
    }
    MPI_Finalize();

    free(A);
    free(Ap);
    return 0;
}

```

Que2.2)

Performance Results with fixed Matrix size with varying Number of Cores

Case 1: 4 cores, N 10000 time spent 2.490000 sec
 Case 2: 8 cores, N 10000 time spent 1.560000 sec
 Case 2: 16 cores, N 10000 time spent 0.930000 sec
 Case 2: 32 cores, N 10000 time spent 0.540000 sec

The scaling of performance with the number of cores is not a linear function. It always actually less than the ratio of the number of the cores, and depends on the workload on all the cores as well as whether the cores are on the same node or not.

Performance Results with varying Matrix size and fixed Number of Core

Case 2: 8 cores, N 100 time spent 0.000-00 sec
 Case 3: 8 cores, N 500 time spent 0.010000 sec
 Case 3: 8 cores, N 1000 time spent 0.020000 sec
 Case 4: 8 cores, N 5000 time spent 0.370000 sec
 Case 4: 8 cores, N 10000 time spent 1.560000 sec

With the increase in the number of elements in a matrix the performance decreases and is slightly higher drop than the ratio of the number of elements in matrix. Like from N= 5000 to N = 10000 the performance decrease is slightly more than 4 times.